



**Red Hat**

# Red Hat OpenShift Development I: Introduction to Containers with Podman

DO188





# Let's meet each other!



Francesco Marchioni - [fmarchio@redhat.com](mailto:fmarchio@redhat.com)  
Red Hat Certified Architect (RHCA)

## A spiral-bound notebook with a white cover and black spiral binding. The cover features the title "COURSE SCHEDULING" in large, bold, black letters. Below the title is a colorful, hand-drawn calendar grid. The grid is divided into sections for different days of the week, each with a unique background color and decorative elements. The sections include: "SATURDAY" (pink background, with a drawing of a girl), "SUNDAY" (blue background, with a drawing of a boy), "MONDAY" (green background, with a drawing of a girl), "TUESDAY" (yellow background, with a drawing of a boy), "WEDNESDAY" (orange background, with a drawing of a girl), "THURSDAY" (purple background, with a drawing of a boy), "FRIDAY" (red background, with a drawing of a girl), "SATURDAY" (pink background, with a drawing of a boy), and "SUNDAY" (blue background, with a drawing of a girl). Each section contains a grid of dates and times, with various course names and numbers written in. The notebook is surrounded by school supplies, including pens, pencils, paper clips, and a bowl of fruit. The background is a wooden surface.

**9:00 AM – 3 PM**



# Online training Attendance



Please remain online during  
Course hours.



# OpenShift Learning Path

# OpenShift Training

DO188  
OpenShift  
Container  
Development I

Podman  
Container  
Images  
Compose

DO288  
OpenShift  
Development II

Build / Deploy/  
Manage  
RHOCP  
Deployments

DO180  
OpenShift  
Admin I

DO280  
OpenShift  
Admin II

Manage/Deploy  
Resources  
Security  
Network  
Operators/Cluster  
update

DO380  
OpenShift  
Admin III

Advanced Security  
Backup/Restore  
Scheduling/Monitor  
Logging  
Gitops

DO480  
OpenShift  
Admin IV

RHACM  
Quay.io  
RHACS

←DevOps→




An abstract graphic on the left side of the slide, rendered in various shades of red. It features a vertical stack of server racks at the bottom, a cloud with a keyhole icon, a large upward-pointing arrow, and several curved arrows suggesting movement or flow. There are also some 'X' and 'O' symbols scattered throughout the design.


# Welcome to World of Containers!



# Problems a containers tries to solve....




# Problems a containers tries to solve.....



# What is a Container ????


# Evolution of Application Deployments



# Benefits and Challenges of using Containers



Let's explore this more in detail!




# Container Concepts and Terminology Chapter 1

# a container is the smallest compute unit

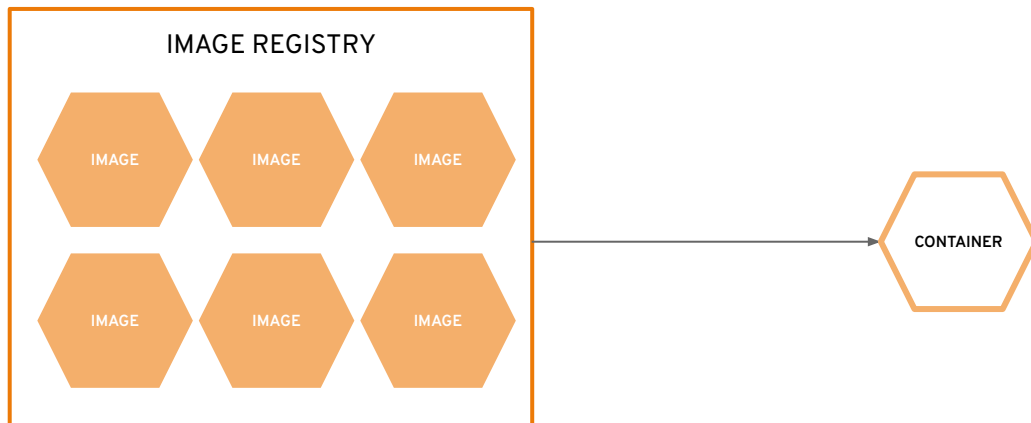


# Containers are created from Container Images

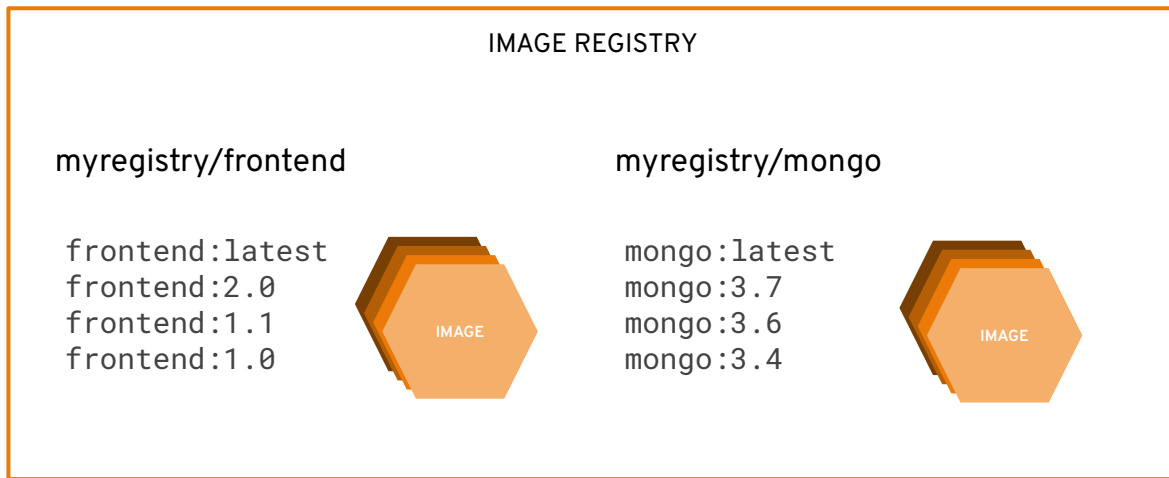




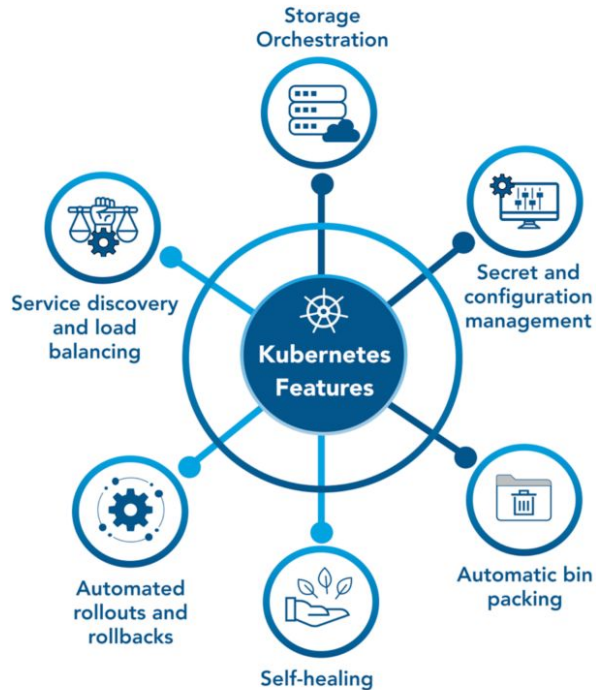
# container images are stored in an image registry




# an Image Repository contains all versions of an image in the image registry



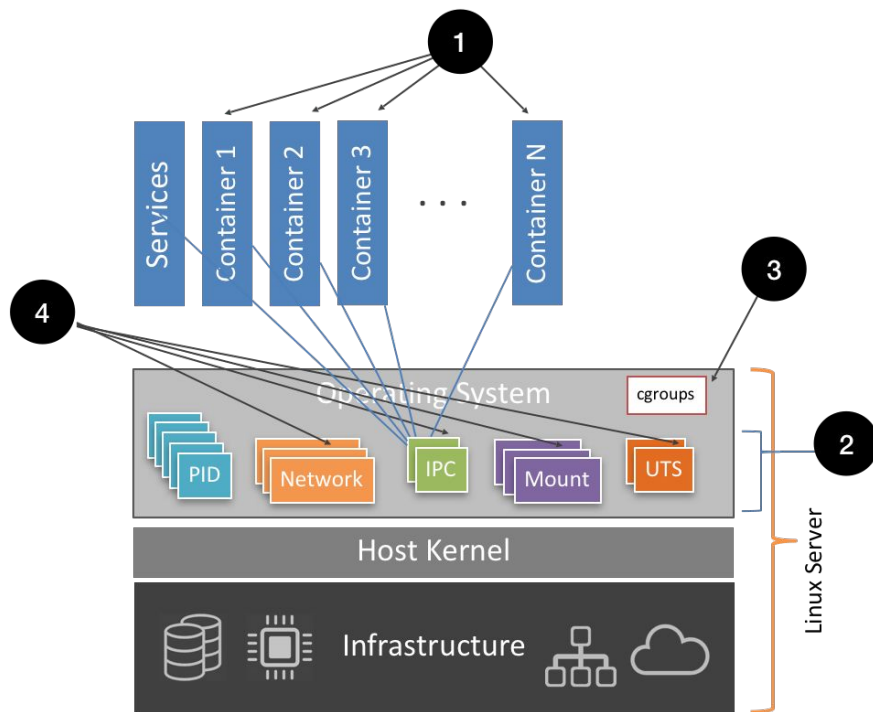
# From Containers to Kubernetes



# Kubernetes Architecture



# Inside Containers



- **(#1)** Each Container has a unique PID namespace running a group of processes.
- **(#2)** Namespaces exist at Kernel level
- **(#3)** Control Groups (cgroups) are used to limit the access to system resources
- **(#4)** Containers will have other dedicated namespaces or will use the default namespace.

# Kubernetes Pod




IP: 10.5.1.20

Pod with single App Container.  
This is the most common and  
preferred type of Pod.



IP: 10.5.1.73

Pod with primary App Container  
and helper sidecar Container  
(i.e. using helper as data change watcher)

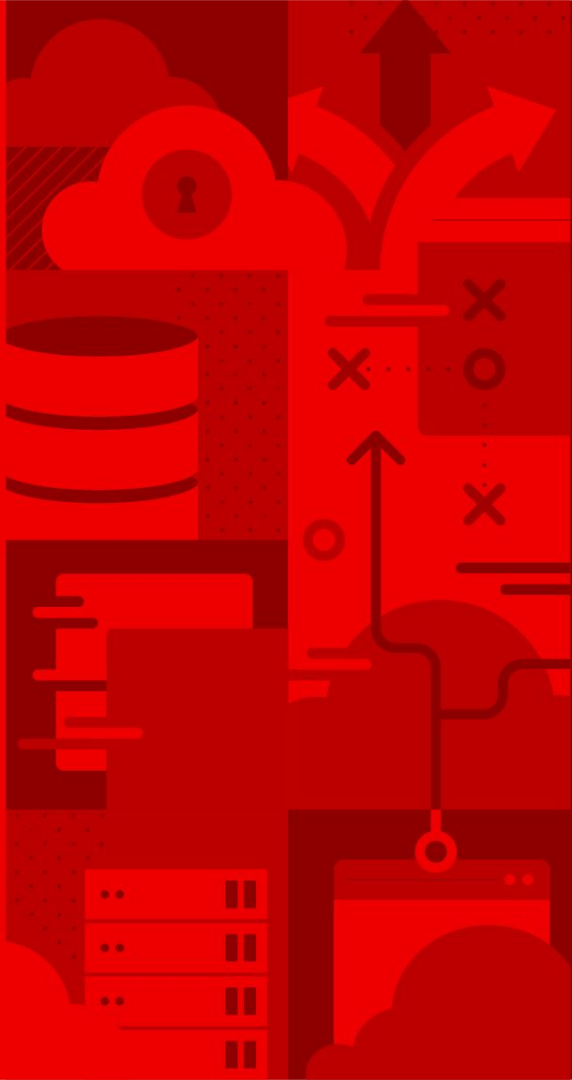


IP: 10.5.1.42

Pod with App Container  
and Sidecar Container  
(i.e. when using Envoy Proxy)

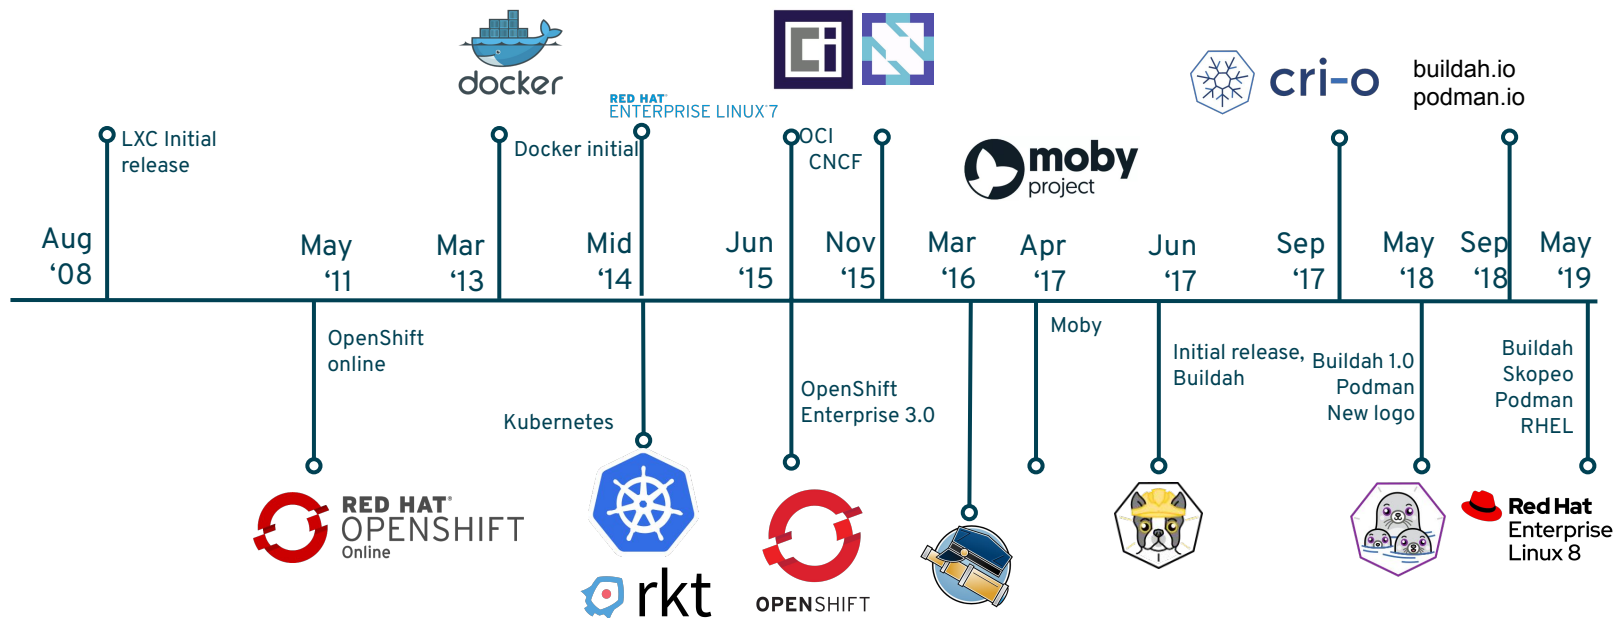


*Container Volumes:* The Volumes can be shared among Containers in the same Pod

An abstract graphic on the left side of the slide, rendered in various shades of red. It features a vertical stack of server racks at the bottom, a cloud with a keyhole icon, a large upward-pointing arrow, and several curved arrows suggesting movement or flow. There are also some 'X' and 'O' symbols scattered throughout the design.

# Creating Containers with Podman Chapter 2

## Container innovation continues





# Introducing Podman and its companion tools

Podman, Buildah and Skopeo



Podman

- the POD MANager - is a tool for managing containers and images, volumes mounted into those containers, and pods made from groups of containers.



Buildah

is a tool that facilitates building Open Container Initiative (OCI) container images without a full container runtime or daemon installed




Skopeo

works with remote images registries - retrieving information, images, signing content

## From Docker to Podman

Helping the transition between technologies



Podman is designed to provide a smooth transition experience to Docker users thanks to CLI compatibility and OCI compliance.

However, knowing the technology foundations and main differences between the two engines is useful to provide a better migration experience.

# Main differences between Docker and Podman

## Overview of the main technology pillars

Docker is a daemon-based container engine that consists of three fundamental pillars:

- Docker daemon, a service running in background that supervises the containers orchestration with a client/Server model
- Docker REST API, an interface to interact with the daemon
- Docker CLI, a CLI interface to manage containers and related resources

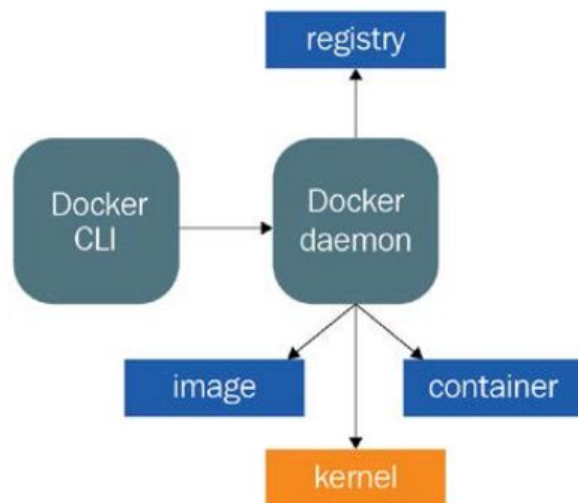
Podman is a container engine that enables users to manage containers, images, as well as storage and network resources:

- Daemonless architecture with fork/exec model
- Standard libpod library for container lifecycle
- Native support for rootless containers and Kubernetes pods
- Full support for OCI runtime/images/storage specs

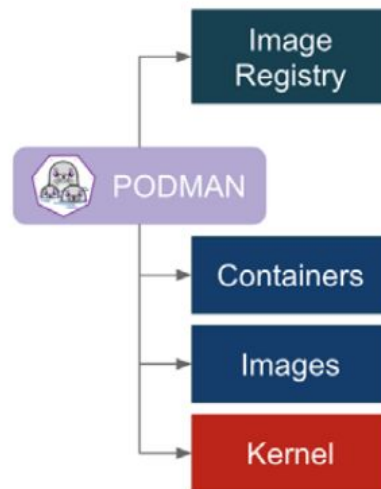
## A 10000 ft view

An high level overview of the two different architectures

### Docker



### Podman



## # sed 's/docker/podman/g'

Different engines, same driving experience

Podman provides a compatible command line interface for most of Docker standardized commands.

Simply swap commands or use podman-docker default aliasing.

```
# Running an Nginx container with Docker
$ docker run -d --rm docker.io/library/nginx
```

```
# Running an Nginx container with Podman
$ podman run -d --rm docker.io/library/nginx
```

# Running containers and pods

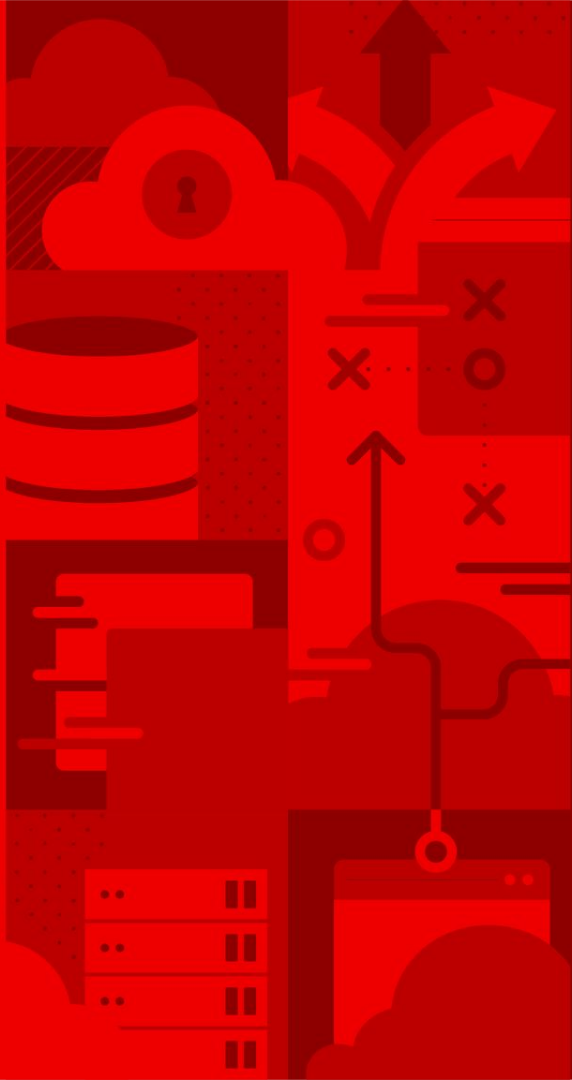
## Pods as minimal execution units

Podman supports the execution of containers and pods. A pod is a single execution unit where multiple containers can share Linux namespaces.

```
# Initializing an empty pod
$ podman pod create --name wp-pod

# Creating containers inside the pod
$ podman create --pod wp-pod --name db -d docker.io/library/mysql
$ podman create --pod wp-pod --name wp -d
  docker.io/library/wordpress


# Running the pod
$ podman pod start wp-pod
```



# Container Networking Chapter 2

# Podman Networking: Introduction

- Each container runs in its own network namespace by default.
- Podman supports user-defined networks to allow communication between containers.



- Network management differs for rootful and rootless containers.
- Podman uses different network backends: CNI (legacy) and Netavark (default since v4).







# Creating and Using Podman Networks

- Use `podman network create <name>` to define a custom bridge network.
- Attach a container to a network with `--net <network>` during `podman run`.
- Containers on the same network can communicate via IP or DNS (if enabled).
- Example:

```
$ podman network create mynet
```





```
$ podman run --net mynet --name web nginx
```

# Effect of Joining a Network

-  Containers in a network (e.g., bridge) can reach each other  
Shared network namespace enables communication via IP or DNS.
-  DNS resolution works if enabled in the network settings  
**Container names are resolvable within the same network.**
-  Containers outside any defined network (e.g., using pasta) are isolated  
**No communication between containers, even via IP.**
-  Network connection affects access from host, container linking, and service discovery

Proper network setup is crucial for multi-container applications.

# Bridge vs Pasta Networking

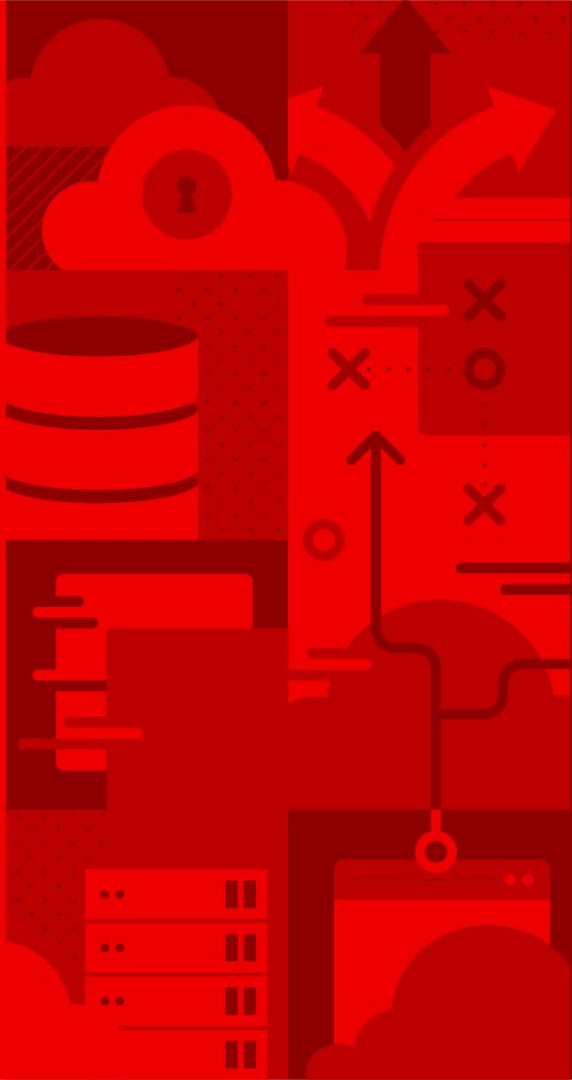
-  **Bridge**: Shared virtual network; supports DNS, IP access, and container communication. Ideal for multi-container apps that need to talk to each other.
-  **Pasta**: Rootless-only; per-container isolation, no inter-container communication. Designed for strong security and minimal configuration.
-  Bridge is default for rootful containers; Pasta is default for rootless (if no `--net`). The default mode depends on user privileges.
-  Containers using pasta cannot be added to bridge networks later. Network modes are not compatible post-creation.



# Summary and Best Practices





- Use **bridge** networks for multi-container applications and service discovery.
- Use **pasta** for secure, isolated rootless containers with minimal config.
- *Always* define networks explicitly in scripts or production setups.
- Inspect networks with ``podman network inspect <name>``.

## Dashboard per lab **basics-exposing**



# Container Registries Chapter 3

# Introduction to Registries


-  Advantages of using Registries
- Container Registries and Podman
- Objectives:
  -  Understand how container images are packaged with all necessary dependencies.
  -  Learn how registries are used to store and share container images.
  -  Get an overview of managing images with Podman.

# What is the Purpose of Registries?

Registries allow you to store and share container images in a controlled way.

Examples of Registries:

- Quay.io
- Red Hat Registry
- Docker Hub
- Amazon ECR



Example Command:

```
podman pull registry.redhat.io/ubi8/ubi:8.6
```

This pulls an image from the Red Hat Registry.



# Red Hat Registries

Red Hat Registries:

- **registry.access.redhat.com**: No authentication
- **registry.redhat.io**: Requires authentication.



Red Hat Ecosystem Catalog:

- Visit [catalog.redhat.com](https://catalog.redhat.com) to search images and view technical details.

Useful Container Images:

- UBI 9: [registry.access.redhat.com/ubi9](https://registry.access.redhat.com/ubi9) – Base image on RHEL 9.
- Python 3.9: [registry.access.redhat.com/ubi9/python-39](https://registry.access.redhat.com/ubi9/python-39) – UBI-based image with Python runtime.
- Node.js 18: [registry.access.redhat.com/ubi9/nodejs-18](https://registry.access.redhat.com/ubi9/nodejs-18) – UBI-based image with Node.js runtime.
- Go Toolset: [registry.access.redhat.com/ubi9/go-toolset](https://registry.access.redhat.com/ubi9/go-toolset) – Contains Go runtime.

# Managing Registries with Podman and Skopeo

## Podman:

- Uses fully qualified image names:
- Format: Registry URL/User or organization/Image repository:Image tag
- If using an unqualified name (e.g., ubi8/python-39), Podman searches registries defined in `/etc/containers/registries.conf`.



## Skopeo:

- A command-line tool to work with container images without local storage.
- Functions:
  - Inspect remote images.
  - Copy images between registries.
  - Sign images using OpenPGP keys.
  - Convert image formats (e.g., Docker to OCI).



# Authentication and Credentials

## Authentication Requirements:

- Some registries (e.g., registry.redhat.io) need user authentication.

## Logging In:

- Use the command: `podman login registry.redhat.io`
- Then enter your username and password.

## Credential Storage:

- Podman stores credentials in `${XDG_RUNTIME_DIR}/containers/auth.json`.



## Best Practices:

- Always use fully qualified image names to avoid pulling duplicate or unintended images. Do not use `latest()` images in production!

---

# Managing Images

## Managing Images with Podman

- Managing Images with Podman
- Objectives:
  - Learn how to pull, tag, build, push, inspect, and remove container images.
  - Understand best practices for image versioning and management.

# Overview of Image Management

## Key Operations:

- Tagging: Map image versions to product updates.
- Pulling: Download images from registries (e.g., using podman pull).
- Building: Create images from a Containerfile with podman build.
- Pushing: Share images by pushing them to remote repositories.
- Inspecting: Retrieve image metadata using podman image inspect.
- Removing: Clean up unused or dangling images to free storage.

# Image Versioning & Tagging

## Deployment Artifacts:

- Container images package software and are versioned similar to product releases.

## Semantic Versioning:

- Format: MAJOR.MINOR.PATCH
- MAJOR: Incompatible changes
- MINOR: Backward-compatible improvements
- PATCH: Bug fixes

## Naming Format:

- [`<repository>/<namespace>/`]`<image name>[:<tag>]`

# Pulling and Building Images

## Pulling Images:

- Use: `podman pull IMAGE_NAME` (e.g., `podman pull quay.io/argoproj/argocd`)
- Note: Unqualified names are resolved using the registries defined in `/etc/containers/registries.conf`.

## Building Images:

- Create images from a Containerfile with:

```
podman build --file Containerfile --tag  
quay.io/YOUR_QUAY_USER/IMAGE_NAME:TAG
```

## Storage Locations:

- Non-root: `~/.local/share/containers`
- Root: `/var/lib/containers`



## Pushing and Inspecting Images

### Pushing Images:

- Authentication: Login with `podman login REGISTRY`
- Push using:

```
podman push quay.io/YOUR_QUAY_USER/IMAGE_NAME:TAG
```

### Inspecting Images:

- Use: `podman image inspect IMAGE_NAME` to view metadata

### Go Templating:

- Extract specific details with the `--format` option, e.g.:
- `podman image inspect IMAGE --format="{{.Config.Cmd}}"`

## Removing Images

- Delete a single image:

```
podman image rm <IMAGE_NAME:TAG>
```

- Force removal (stops dependent containers):

```
podman image rm -f IMAGE_NAME:TAG
```

- Remove a Container

```
podman rm <CONTAINER>
```

- Remove all dangling images [\*]:

```
podman image prune -a
```

[\*] = Untagged images generated during the build and not used any more

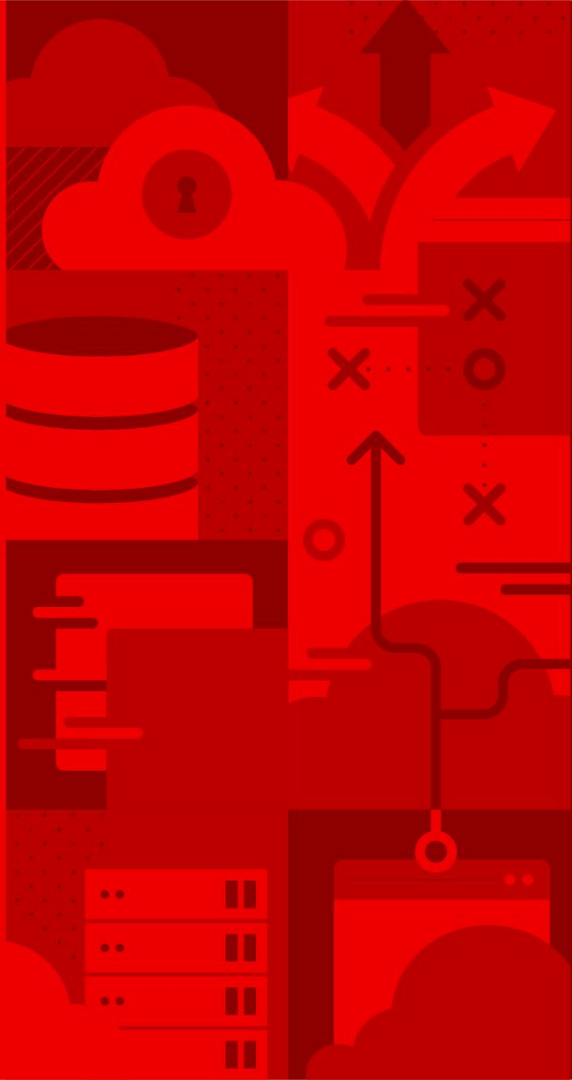
## Exporting and Importing

- Export: Save a container's filesystem as a .tar file:

```
podman export -o mytarfile.tar CONTAINER_ID
```

- Import: Create an image from an exported .tar file:

```
podman import mytarfile.tar IMAGE_NAME:TAG
```



# Custom Container Images Chapter 4

# Introduction

## What is a Containerfile?

- A list of instructions for building a container image.
- Each instruction creates a new image layer.
- Similar to Dockerfiles but independent of any container runtime engine.
- Used to create multiple containers from the same image.

## Why Use Containerfiles?

- Standardized way to define and build container images.
- Compatible with tools like Podman and Docker.
- Enables version control and reproducibility.

# Exporting and Importing

## What is a Base Image?

- The starting point for building a container image.
- Defines the Linux distribution, package manager, and preinstalled dependencies.

## Types of Red Hat Universal Base Images (UBI):

- Standard – Includes DNF, systemd, gzip, and tar.
- Init – Uses systemd to manage multiple applications in one container.
- Minimal – A smaller image with the microdnf package manager.
- Micro – The smallest image with only essential components.

## Where to Find UBIs?

- Available in the Red Hat Container Catalog.

# Exporting and Importing

- Basic Instructions:
  - **FROM** – Defines the base image.
  - **WORKDIR** – Sets the working directory inside the container.
  - **COPY / ADD** – Copies files into the container image.
  - **RUN** – Executes commands during image build.
- Execution and Configuration:
  - **ENTRYPOINT** – Specifies the executable for the container.
  - **CMD** – Provides arguments for the **ENTRYPOINT**.
  - **USER** – Changes the user inside the container.
- Metadata and Networking:
  - **LABEL** – Adds metadata.
  - **EXPOSE** – Documents ports used by the container.
  - **ENV** – Defines environment variables.

## Example Containerfile

```
# Apache Web Server Containerfile

FROM registry.redhat.io/ubi8/ubi:8.6

LABEL description="Custom httpd container image"

RUN yum install -y httpd

EXPOSE 80

ENV LogLevel "info"

ADD http://someserver.com/filename.pdf /var/www/html

COPY ./src/ /var/www/html/

USER apache

ENTRYPOINT ["/usr/sbin/httpd"]

CMD ["-D", "FOREGROUND"]
```

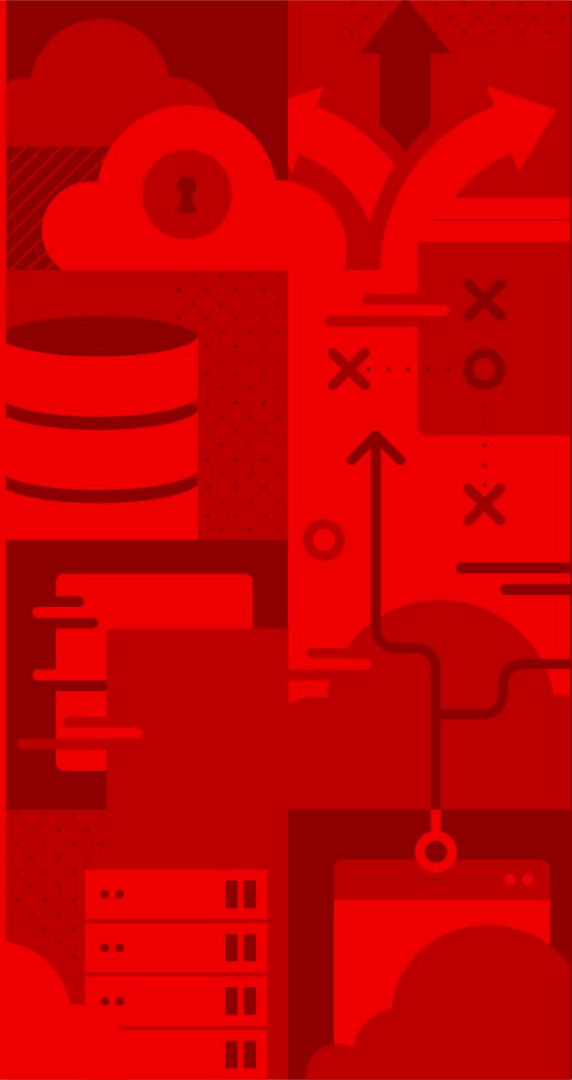
### Key Instructions in This Example:

- Uses `ubi8/ubi:8.6` as base image.
- Installs Apache (`httpd`) using `RUN`.
- Defines `EXPOSE` for port 80.
- Adds files with `ADD` and `COPY`.
- Runs as user `apache` for security.
- Uses `ENTRYPOINT` and `CMD` to start the web server.



# Images Tagging Best Practices

- Tagging Images:
  - Format: `image-name:tag`
  - Example: `my-app:1.0`
  - Default tag is `latest` if omitted.
- Best Practices:
  - Always specify a tag to prevent unintended updates.
  - Use meaningful labels (`LABEL`) for better organization.
  - Avoid running containers as root (`USER`).
  - Minimize image size by using lightweight base images.
- Cleaning Up Unused Images:
  - `podman image prune` – Removes unused images.
  - `podman rmi <image>` – Deletes a specific image.



# Rootless Containers Chapter 4

# Introduction to Rootless Podman

What is Rootless Podman?

- A container runtime that allows running containers without requiring root privileges.
- Enhances security by isolating containerized applications from the host system.
- Uses user namespaces to map container users to non-root host users.



## Why use Rootless Podman ?

- **Security Benefits:**
  - Reduces the risk of privilege escalation.
  - Limits potential attack surfaces in case of container escape.
- **Flexibility:**
  - Enables non-root users to run containers without admin intervention.
  - Works well in multi-user environments.
- **Compatibility:**
  - Supports standard container workloads while ensuring security.

## How Rootless Podman works

- **Key Concepts:**
  - Podman starts containers as a user-space process without requiring root access.
  - Uses user namespaces to map container UIDs/GIDs to unprivileged host users.
  - The container runtime does not need elevated privileges.
- **Process Flow:**
  - Podman starts the container.
  - The containerized process attaches to the systemd parent process.

# Prerequisites for Running Rootless Containers

- **Kernel Features Required:**
  - cgroup v2: Allows resource control without root privileges.
  - [slirp4netns](#): Enables user-mode networking for rootless containers.
  - [fuse-overlayfs](#): Manages the Copy-On-Write (COW) filesystem efficiently.
- **Setting Up Subordinate User IDs:**
  - `/etc/subuid` and `/etc/subgid` define UID/GID mappings.
  - Use `usermod --add-subuids` and `podman system migrate` to configure.

## Changing the Container User

- Default Behavior:
  - Containers typically run as **root** inside, posing security risks.
  - Exploits can escalate privileges from the container to the host.
- Best Practice:
  - Use the **USER** directive in the Containerfile:

```
FROM registry.access.redhat.com/ubi9/ubi
```

```
RUN adduser --no-create-home --system --shell /usr/sbin/nologin appuser
```

```
USER appuser
```

```
CMD ["python3", "-m", "http.server"]
```

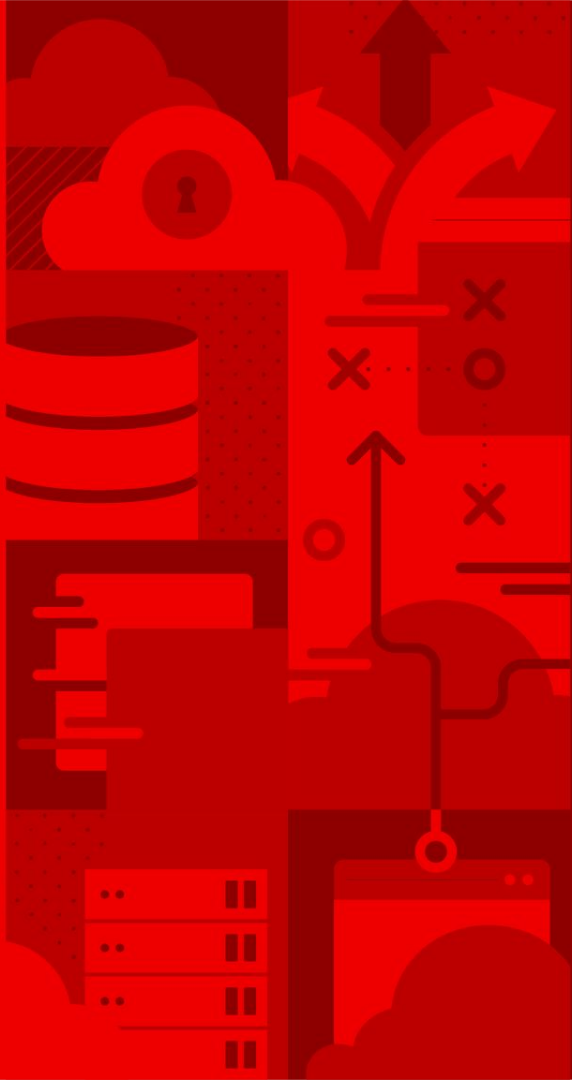
# User Mapping in Rootless Containers

- How User Mapping Works:
  - Container users map to unprivileged host users.
  - Example mapping in `/etc/subuid`:  
student:100000:65536
  - `podman top <container_id> huser user` shows mapped users.
- Example:
  - Root inside the container maps to UID 1000 on the host.



## Limitations of Rootless Containers

- Common Restrictions:
  - Cannot bind to privileged ports (e.g., 80, 443) unless configured (`sysctl`).
  - Some applications requiring `root` (e.g., ping) may need workarounds.
  - Not all applications can be containerized without root.
- Workarounds:
  - Use port forwarding instead of binding directly to privileged ports.
  - Adjust kernel parameters to allow specific capabilities (e.g., `ping`)

An abstract graphic on the left side of the slide, rendered in various shades of red. It features a vertical stack of server racks at the bottom, a cloud with a keyhole icon, a large upward-pointing arrow, and several curved arrows indicating data flow or movement. There are also some 'X' and 'O' symbols scattered throughout the design.

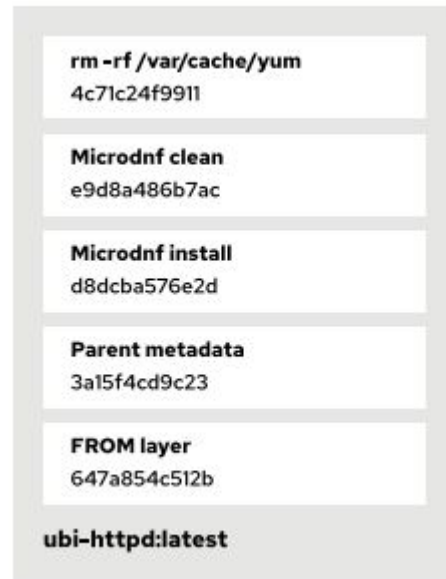
# Managing Storage Chapter 5

# Introduction to Volume Mounting in Podman

- Objective: Understand the process of mounting volumes and its common use cases.
- **Why Use Volume Mounting?**
  - Persistent storage across container deletions
  - Efficient data sharing between containers
  - Performance improvements for write-heavy applications

# Understanding the Copy-on-Write (COW) File System

- Container Image Layers
  - Each modification creates a new read-only layer.
  - The final image is a union of multiple layers.
- Example Image Structure (using **podman image tree**)
  - Multiple immutable layers form the base of a container.
- Implication: Containers share base layers, but each container has a separate read/write layer.



## The Need for External Storage

- **Challenges with COW File System:**
  - Write operations are inefficient due to layer copying.
  - Data is lost when the container is removed.
- **Solution:** Store data outside of the container using bind mounts or volumes.

# Types of External Mounts in Podman

## 1. Bind Mounts

- Maps a host directory into the container.
- Requires manual management of permissions.

## 2. Volumes

- Managed by Podman.
- More portable and persistent.

## 3. Tmpfs Mounts

- Memory-backed, ephemeral storage.

# Using Bind Mounts

## Command Syntax:

```
podman run -p 8080:8080 --volume /www:/var/www/html:ro  
registry.access.redhat.com/ubi8/httpd-24:latest
```

- Key Considerations:
  - Read-only (**ro**) prevents modifications.
  - Ensure correct file permissions.

# Troubleshooting Bind Mounts

- **File Permission Issues:**
  - Use `podman unshare ls -l /www/` to check permissions.
- **SELinux Issues:**
  - Check with `ls -Zd /www`

Fix using `:z` (shared access) or `:Z` (exclusive access):

```
podman run -p 8080:8080 --volume /www:/var/www/html:Z
registry.access.redhat.com/ubi8/httpd-24:latest
```



# Using Podman Volumes

- Why Use Volumes?
  - More portable and easier to manage.
  - Ensures consistency across different hosts.
- Creating and Using Volumes:

```
podman volume create http-data  
  
podman run -p 8080:8080 --volume http-data:/var/www/html \  
registry.access.redhat.com/ubi8/httpd-24:latest
```

- How to change mount point volume path in podman?  
<https://access.redhat.com/solutions/7020076>

## Advanced Volume Creation

- Podman volume create allows to use advanced options.
- For example, you can map an NFS volume by setting the appropriate type and device settings [1]

Creating and Using Volumes:

```
$ podman volume create --driver local --opt type=nfs --opt  
o=addr=my-nfs-share.example.com,rw,context="system_u:object_r:container_file_t:s0"  
--opt device=:/opt/container-nfs-storage nfs-volume-01
```

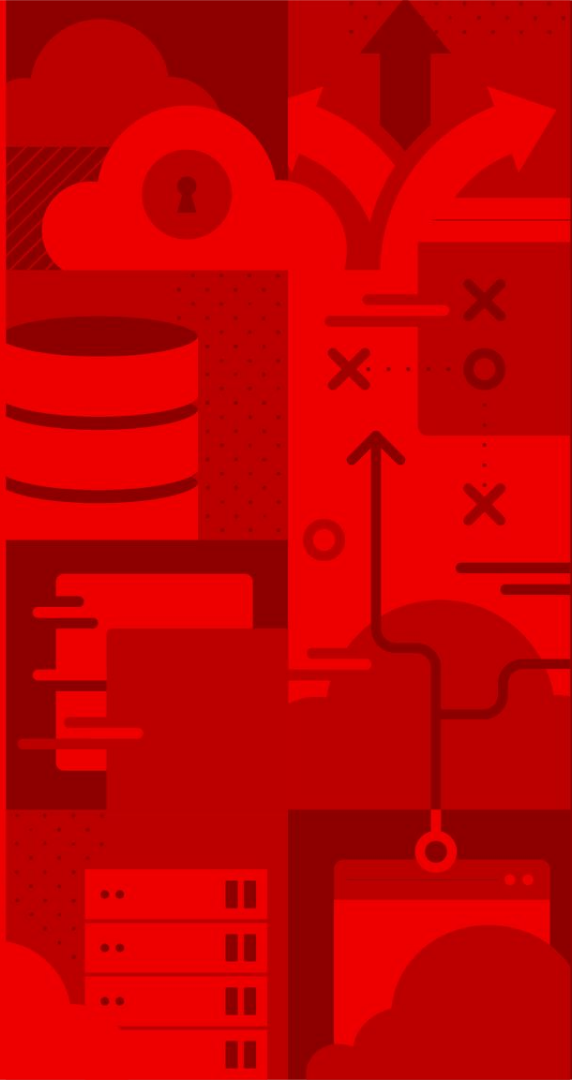
[1] <https://access.redhat.com/solutions/6415411>

# Managing Volumes

- Inspect Volume Details:  
`podman volume inspect http-data`
- Listing All Volumes:  
`podman volume ls`
- Removing Unused Volumes:  
`podman volume rm http-data`

## Summary & Best Practices

- Use Bind Mounts for:
  - Temporary data.
  - Environment-specific configurations.
- Use Volumes for:
  - Persistent data storage.
  - Portability across environments.
  - Better management by Podman.
- Ensure correct permissions when using bind mounts.



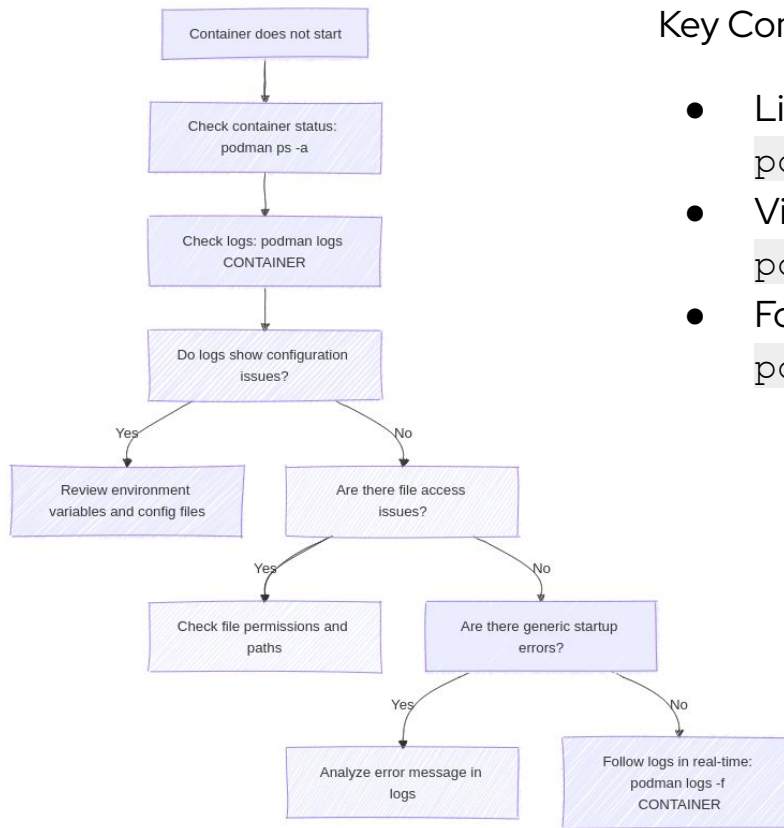
# Troubleshooting Containers Chapter 6

## Summary & Best Practices

### Objectives:

- Read container logs.
- Troubleshoot common container problems.
- Check Red Hat KCS for Podman [Solutions](#)

# Troubleshoot common issues



## Key Commands:

- List all containers:  
`podman ps -a`
- View container logs:  
`podman logs CONTAINER`
- Follow logs in real-time:  
`podman logs -f CONTAINER`

## Podman internal Logs

- How to check Podman's own logs?

```
$ journalctl _COMM=podman
```

```
Nov 19 11:41:08 r95jbu01 podman[12609]: 2024-11-19 11:41:08.159570184 +0900 JST  
m=+0.203488520 system refresh
```

```
Nov 19 11:41:21 r95jbu01 podman[12609]: 2024-11-19 11:41:21.880379337 +0900 JST  
m=+13.924297681 image pull
```

```
02141a49ee4eaf2aa824472c45a281be14e3e85299e117891da1153a5eddc4c7 registry.access.r>
```

```
Nov 19 11:41:36 r95jbu01 podman[12699]: 2024-11-19 11:41:36.476350952 +0900 JST  
m=+0.073530605 image pull
```

```
02141a49ee4eaf2aa824472c45a281be14e3e85299e117891da1153a5eddc4c7 ubi8
```

```
Nov 19 11:41:36 r95jbu01 podman[12699]: 2024-11-19 11:41:36.586357214 +0900 JST  
m=+0.183536868 container create
```

- Use the `--log-level <level>` option to set the logging level.



# Troubleshoot containers networking

## Common Issues:

- Incorrect port mapping.
- No network access between containers.
- DNS resolution failures.



How to troubleshoot containers using nsenter:

<https://access.redhat.com/solutions/1611883>

## Port Mapping Verification:

```
podman port CONTAINER
```

## Check open ports in the container:

```
podman exec -it CONTAINER ss -pant
```

## Use nsenter for network namespace troubleshooting:

```
podman inspect CONTAINER --format '{{.State.Pid}}'
```

```
sudo nsenter -n -t CONTAINER_PID ss -pant
```

## Troubleshoot containers connectivity

Verify Network Assignment:

```
podman inspect CONTAINER --format='{{.NetworkSettings.Networks}}'
```

Check if DNS is enabled:

```
podman network inspect NETWORK
```

## Podman Events for troubleshooting



- Monitor container events:

```
podman events --stream=false
```

- Filter events by type:

```
podman events --filter event=create --filter type=container  
--stream=false
```

- View past events:

```
podman events --since 5m --stream=false
```

# Red Hat OpenShift and Kubernetes Core Concepts




# A container is the smallest compute unit




CONTAINER

## Everything runs in pods




Containers are wrapped in pods which are units of deployment and management



POD


10.140.4.44



POD


10.15.6.55

ReplicationControllers & ReplicaSets ensure a specified number of pods are running at any given time

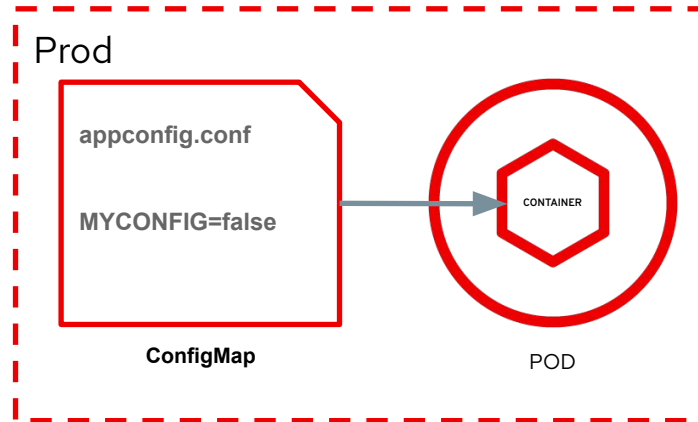
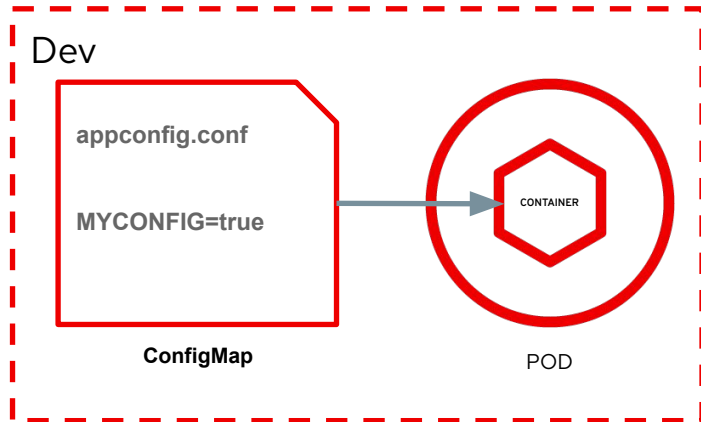




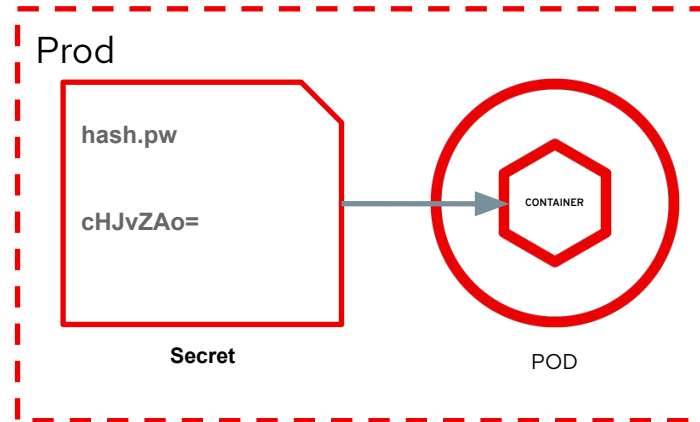
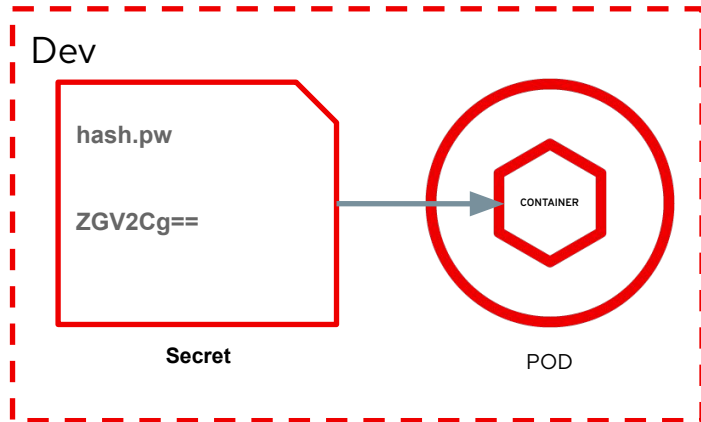
## Deployments define how to roll out new versions of Pods



## Configmaps allow you to decouple configuration artifacts from image content

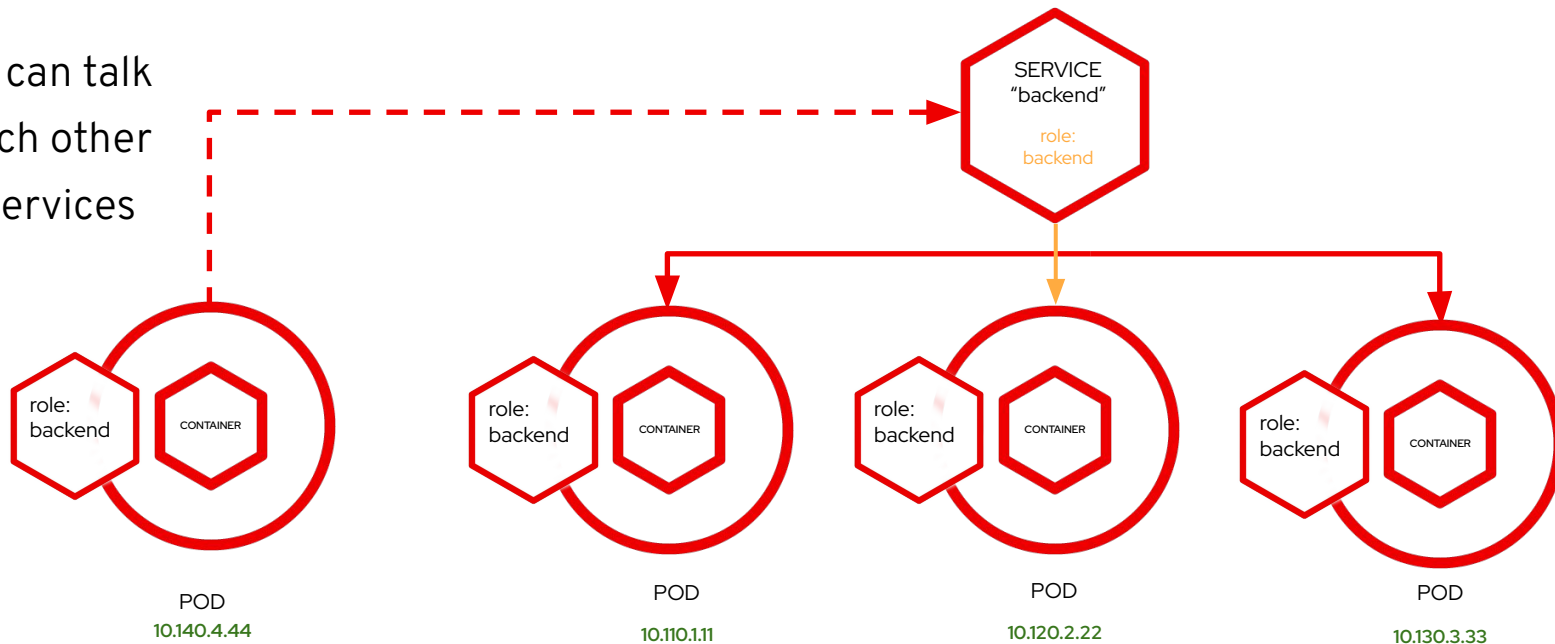


Secrets provide a mechanism to hold sensitive information such as passwords




## Services provide internal load-balancing and service discovery across pods

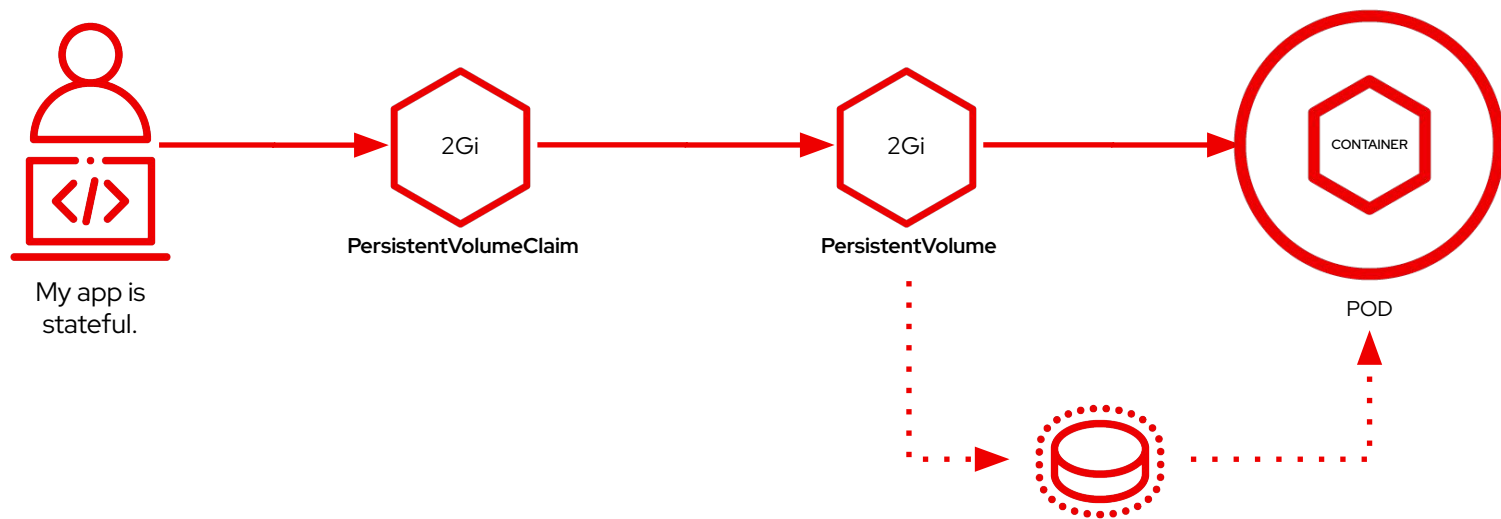
Apps can talk  
to each other  
via services




## Routes make services accessible to clients outside the environment via real-world urls



## Persistent Volume and Claims



## Liveness and Readiness





# Final review

## Chapter 9



## Architecture Overview

